



Lab meeting

Deep dive into deep RL

Andrea Pierré

February 12th, 2024

Brown University

Outline


1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

Outline

1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

How do neurons in the LEC integrate sensory and spatial information?

- **Piriform Cortex** encodes olfactory information
- **Hippocampus** encodes spatial information
- **Lateral Entorhinal Cortex (LEC)** encodes both olfactory & spatial information



Why modeling?

- Having a reliable model to make predictions
- Simulation may uncover insight's from the real data
- Derive principles from the simulation which can be checked in the experiment


Why modeling?

- Having a reliable model to make predictions
- Simulation may uncover insight's from the real data
- Derive principles from the simulation which can be checked in the experiment

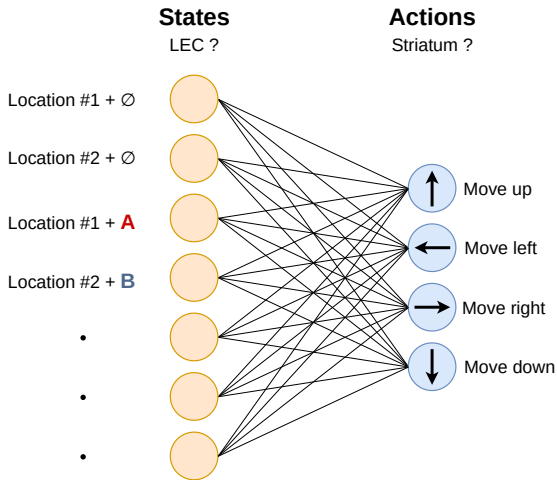
Why modeling?

- Having a reliable model to make predictions
- Simulation may uncover insight's from the real data
- Derive principles from the simulation which can be checked in the experiment

The half-triangle task



Mapping states to action




Outline

1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions


Deep RL lessons learned

- Deep RL is different from supervised learning \rightarrow the data to optimize on is not fixed (moving target)
- DQN tricks:




Deep RL lessons learned

- Deep RL is different from supervised learning \rightarrow the data to optimize on is not fixed (moving target)
- DQN tricks:
 - Replay buffer \rightarrow save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - Exploration warm up with ϵ -greedy \rightarrow experience more diverse data
 - Batching \rightarrow update the weights of the network based on several data examples at the same time instead of only one
 - Target network \rightarrow use 2 networks to stabilize learning (one of them is updated with a lag)




Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - Replay buffer → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - Exploration warm up with ϵ -greedy → experience more diverse data
 - Batching → update the weights of the network based on several data examples at the same time instead of only one
 - Target network → use 2 networks to stabilize learning (one of them is updated with a lag)




Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - **Exploration warm up** with ϵ -greedy → experience more diverse data
 - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
 - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - **Exploration warm up** with ϵ -greedy → experience more diverse data
 - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
 - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - **Exploration warm up** with ϵ -greedy → experience more diverse data
 - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
 - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - **Exploration warm up** with ϵ -greedy → experience more diverse data
 - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
 - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - **Exploration warm up** with ϵ -greedy → experience more diverse data
 - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
 - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)



Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
 - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
 - **Exploration warm up** with ϵ -greedy → experience more diverse data
 - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
 - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)



Deep RL implementation

Algorithm 1: Deep Q-Network (DQN)

Initialize replay memory D to capacity N

Initialize action-value network Q with random weights θ

Initialize target action-value network \hat{Q} with random weights $\theta^- = \theta$

for $episode \leftarrow 1 \dots M$ do

$state \leftarrow reset(env)$

$done \leftarrow False$

 while $done \neq True$ do

$Q \leftarrow forward_pass(state)$

 /* 4 action values vector */

$action \leftarrow \epsilon_{greedy}(action_space, state, Q)$

$state_{new}, reward, done \leftarrow env.step(action, state)$

 Store transition (state, action, reward, next_state, done) in D

 Sample random minibatch of transitions from D

$Q \leftarrow forward_pass(state_{new})$

 /* 4 action values vector */

$Q_{new} \leftarrow reward + \gamma \max(\hat{Q})$

 /* scalar */

$y \leftarrow \max(Q)$

 /* scalar */

 if $done = True$ then

$\hat{y}_{pred} \leftarrow reward$

 /* scalar */

 else

$\hat{y}_{pred} \leftarrow Q_{new}$

 /* scalar */

 end

$Loss \leftarrow (y - \hat{y}_{pred})^2$

 Perform a gradient descent step on Loss with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

 end

end

Outline


1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

How deep RL feels like









Networks architectures










$28 \rightarrow 54 \rightarrow 54 \rightarrow 54 \rightarrow 4$




Rewards & steps



Loss, rewards & steps distributions



Policy learned (when it works)



Outline

1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

Adding memory into the task


Current environment

step	location	cue	reward
1	2	No odor	0
2	3	No odor	0
3	4	Odor A	0
4	3	Odor A	0
5	2	Odor A	0
6	1	Odor A	0
7	0	Odor A	10

With memorization needed


step	location	cue	reward
1	2	\emptyset	0
2	3	\emptyset	0
3	4	Odor A	0
4	3	\emptyset	0
5	2	\emptyset	0
6	1	\emptyset	0
7	0	\emptyset	10

RNN for memorization and sequence modeling




Feedback connectivity


A.



B.




C.





Network of networks

(a) Single-region RNN



(b) Multi-region RNN (mRNN)



Questions ?