



# Lab meeting

Deep dive into deep RL

---

Andrea Pierré

February 12<sup>th</sup>, 2024

Brown University

# Outline


1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

# Outline

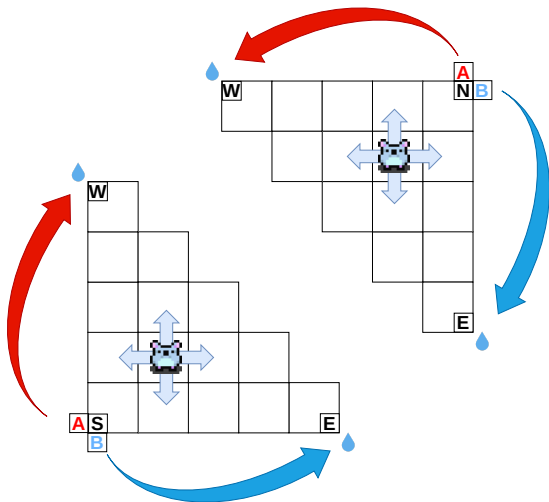
1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

# How do neurons in the LEC integrate sensory and spatial information?


- **Piriform Cortex** encodes olfactory information
- **Hippocampus** encodes spatial information
- **Lateral Entorhinal Cortex (LEC)** encodes both olfactory & spatial information



# The half-triangle task



# Mapping states to action




# Outline

1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

# Deep RL lessons learned


- Deep RL is different from supervised learning  $\rightarrow$  the data to optimize on is not fixed (moving target)
- DQN tricks:






# Deep RL lessons learned

- Deep RL is different from supervised learning  $\rightarrow$  the data to optimize on is not fixed (moving target)
- DQN tricks:
  - Replay buffer  $\rightarrow$  save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - Exploration warm up with  $\epsilon$ -greedy  $\rightarrow$  experience more diverse data
  - Batching  $\rightarrow$  update the weights of the network based on several data examples at the same time instead of only one
  - Target network  $\rightarrow$  use 2 networks to stabilize learning (one of them is updated with a lag)




# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - Replay buffer → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - Exploration warm up with  $\epsilon$ -greedy → experience more diverse data
  - Batching → update the weights of the network based on several data examples at the same time instead of only one
  - Target network → use 2 networks to stabilize learning (one of them is updated with a lag)




# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - **Exploration warm up** with  $\epsilon$ -greedy → experience more diverse data
  - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
  - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - **Exploration warm up** with  $\epsilon$ -greedy → experience more diverse data
  - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
  - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - **Exploration warm up** with  $\epsilon$ -greedy → experience more diverse data
  - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
  - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - **Exploration warm up** with  $\epsilon$ -greedy → experience more diverse data
  - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
  - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)




# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - **Exploration warm up** with  $\epsilon$ -greedy → experience more diverse data
  - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
  - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)



# Deep RL lessons learned

- Deep RL is different from supervised learning → the data to optimize on is not fixed (moving target)
- DQN tricks:
  - **Replay buffer** → save the data experienced in a buffer and sample from it to break the temporal correlation of the data
  - **Exploration warm up** with  $\epsilon$ -greedy → experience more diverse data
  - **Batching** → update the weights of the network based on several data examples at the same time instead of only one
  - **Target network** → use 2 networks to stabilize learning (one of them is updated with a lag)





# Deep RL implementation

---

## Algorithm 1: Deep Q-Network (DQN)

---

Initialize replay memory  $D$  to capacity  $N$

Initialize action-value network  $Q$  with random weights  $\theta$

Initialize target action-value network  $\hat{Q}$  with random weights  $\theta^- = \theta$

for  $episode \leftarrow 1 \dots M$  do

$state \leftarrow reset(env)$

$done \leftarrow False$

    while  $done \neq True$  do

$Q \leftarrow forward\_pass(state)$

        /\* 4 action values vector \*/

$action \leftarrow \epsilon_{greedy}(action\_space, state, Q)$

$state_{new}, reward, done \leftarrow env.step(action, state)$

        Store transition (state, action, reward, next\_state, done) in  $D$

        Sample random minibatch of transitions from  $D$

$Q \leftarrow forward\_pass(state_{new})$

        /\* 4 action values vector \*/

$Q_{new} \leftarrow reward + \gamma \max(\hat{Q})$

        /\* scalar \*/

$y \leftarrow \max(Q)$

        /\* scalar \*/

        if  $done = True$  then

$\hat{y}_{pred} \leftarrow reward$

        /\* scalar \*/

        else

$\hat{y}_{pred} \leftarrow Q_{new}$

        /\* scalar \*/

        end

$Loss \leftarrow (y - \hat{y}_{pred})^2$

        Perform a gradient descent step on Loss with respect to the network parameters  $\theta$

        Every  $C$  steps reset  $\hat{Q} = Q$

    end

end

# Outline


1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions

# How deep RL feels like









# Networks architectures










$28 \rightarrow 54 \rightarrow 54 \rightarrow 54 \rightarrow 4$




# Rewards & steps



# Loss, rewards & steps distributions



# Policy learned (when it works)



# Outline

1. Context
2. Deep RL algorithm & building blocs
3. Results
4. Future directions



# Adding memory into the task


Current environment

step	location	cue	reward
1	2	No odor	0
2	3	No odor	0
3	4	Odor A	0
4	3	Odor A	0
5	2	Odor A	0
6	1	Odor A	0
7	0	Odor A	10

With memorization needed


step	location	cue	reward
1	2	∅	0
2	3	∅	0
3	4	Odor A	0
4	3	∅	0
5	2	∅	0
6	1	∅	0
7	0	∅	10

# RNN for memorization and sequence modeling




# Feedback connectivity


A.



B.




C.




# Network of networks

(a) Single-region RNN



(b) Multi-region RNN (mRNN)



Questions ?