



Research plan





Fleischmann – Nassar joint meeting

Andrea Pierré

May 17, 2024

Brown University

Outline

1. Hypothesis & directions 
2. Experiments & expected results  
3. Roadmap 

Outline

1. Hypothesis & directions 

2. Experiments & expected results  

3. Roadmap 

What do we want to know?

- Understand what the network learns → What **function** does it learn?
- How the constraints of the task affect learning & the representations learned?
- Does the network learn something related to the real neurons? (million \$\$\$ question)

What do we want to know?

- Understand what the network learns → What **function** does it learn?
- How the constraints of the task affect learning & the representations learned?
- Does the network learn something related to the real neurons? (million \$\$\$ question)

What do we want to know?

- Understand what the network learns → What **function** does it learns?
- How the constraints of the task affect learning & the representations learned?
- Does the network learn something related to the real neurons? (million \$\$\$ question)

Compositional hypothesis

- Does the network learn a generalizable policy?
→ How to test it?
 - From indexed locations (current) to coordinate system
 - Merged actions space


Compositional hypothesis

- Does the network learn a generalizable policy?
 - How to test it?
 - From indexed locations (current) to coordinate system
 - Merged actions space


Compositional hypothesis

- Does the network learn a generalizable policy?
 - How to test it?
 - From indexed locations (current) to coordinate system
 - Merged actions space

Compositional hypothesis



Implementation



Outline

1. Hypothesis & directions ?

2. Experiments & expected results  

3. Roadmap 

1) How training impacts the representations learned?

- Feed both coordinates information (Cartesian & polar) to the input layer (+ merge actions spaces in a common one)
- Train on left/right task → we expect the weights are close to zero on Cartesian representation?
- Train on east/west task → we expect the weights are close to zero on polar representation?
- Not clear to me how to extract/define which neurons belong/contribute to Cartesian/polar representation

1) How training impacts the representations learned?

- Feed both coordinates information (Cartesian & polar) to the input layer (+ merge actions spaces in a common one)
- Train on left/right task → we expect the weights are close to zero on Cartesian representation?
- Train on east/west task → we expect the weights are close to zero on polar representation?
- Not clear to me how to extract/define which neurons belong/contribute to Cartesian/polar representation



1) How training impacts the representations learned?

- Feed both coordinates information (Cartesian & polar) to the input layer (+ merge actions spaces in a common one)
- Train on left/right task → we expect the weights are close to zero on Cartesian representation?
- Train on east/west task → we expect the weights are close to zero on polar representation?
- Not clear to me how to extract/define which neurons belong/contribute to Cartesian/polar representation

1) How training impacts the representations learned?

- Feed both coordinates information (Cartesian & polar) to the input layer (+ merge actions spaces in a common one)
- Train on left/right task → we expect the weights are close to zero on Cartesian representation?
- Train on east/west task → we expect the weights are close to zero on polar representation?
- Not clear to me how to extract/define which neurons belong/contribute to Cartesian/polar representation

1) How training impacts the representations learned?



2) Does the network learn a coordinate system?

1. After training, move the population of agents in a translated coordinate system → we expect the population of agents to be able to solve the task with zero shot learning
2. Train with both coordinates information (Cartesian & polar), after training feed incorrect polar angles
 - On the left/right task → we expect the population of agents still solves the task consistently
 - On the east/west task → we expect the network won't converge to a stable policy (i.e all the agents don't solve the task consistently)

2) Does the network learn a coordinate system?

1. After training, move the population of agents in a translated coordinate system → we expect the population of agents to be able to solve the task with zero shot learning
2. Train with both coordinates information (Cartesian & polar), after training feed incorrect polar angles
 - On the left/right task → we expect the population of agents still solves the task consistently
 - On the east/west task → we expect the network won't converge to a stable policy (i.e all the agents don't solve the task consistently)



2) Does the network learn a coordinate system?

1. After training, move the population of agents in a translated coordinate system → we expect the population of agents to be able to solve the task with zero shot learning
2. Train with both coordinates information (Cartesian & polar), after training feed incorrect polar angles
 - On the left/right task → we expect the population of agents still solves the task consistently
 - On the east/west task → we expect the network won't converge to a stable policy (i.e all the agents don't solve the task consistently)

2) Does the network learn a coordinate system?


1. After training, move the population of agents in a translated coordinate system → we expect the population of agents to be able to solve the task with zero shot learning
2. Train with both coordinates information (Cartesian & polar), after training feed incorrect polar angles
 - On the left/right task → we expect the population of agents still solves the task consistently
 - On the east/west task → we expect the network won't converge to a stable policy (i.e all the agents don't solve the task consistently)

2) Does the network learn a coordinate system?




3) Conjunctive odor-place coding

- Train a population of agents, then after training, flip odor A and odor B in the task
- In general we expect to find a population of conjunctive neurons that get active with the combination of both odor and specific location
- Do the conjunctive cells get conserved or remapped?
(Not clear to me, I'd expect they get remapped)



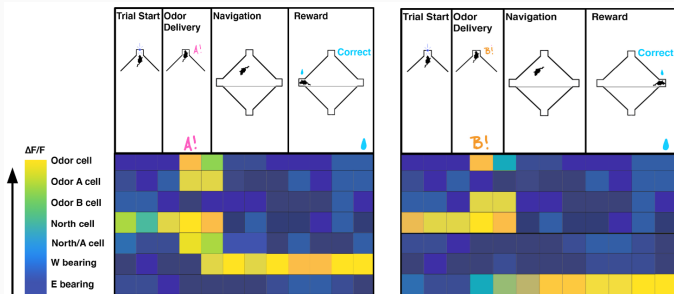
3) Conjunctive odor-place coding

- Train a population of agents, then after training, flip odor A and odor B in the task
- In general we expect to find a population of conjunctive neurons that get active with the combination of both odor and specific location
- Do the conjunctive cells get conserved or remapped?
(Not clear to me, I'd expect they get remapped)



3) Conjunctive odor-place coding


- Train a population of agents, then after training, flip odor A and odor B in the task
- In general we expect to find a population of conjunctive neurons that get active with the combination of both odor and specific location
- Do the conjunctive cells get conserved or remapped?
(Not clear to me, I'd expect they get remapped)



Outline

1. Hypothesis & directions ?

2. Experiments & expected results  

3. Roadmap 

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆

1.1 Code logic for new environment [~1 week]

1.2 Check everything works as expected (unit testing)
[~1 week]

1.3 Bugs? [~1 week]

2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★★⚠

[1 week – 1 month]

3. Experiments

3.1 Task code ★☆☆ [~1 week]

3.2 Training ★★☆☆⚠ [~2 week]

3.3 Analysis code ★★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆

1.1 Code logic for new environment [~1 week]

1.2 Check everything works as expected (unit testing)
[~1 week]

1.3 Bugs? [~1 week]

2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★☆☆⚠

[1 week – 1 month]

3. Experiments

3.1 Task code ★☆☆ [~1 week]

3.2 Training ★★☆☆⚠ [~2 week]

3.3 Analysis code ★★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆

1.1 Code logic for new environment [~1 week]

1.2 Check everything works as expected (unit testing)
[~1 week]

1.3 Bugs? [~1 week]

2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★☆☆⚠

[1 week – 1 month]

3. Experiments

3.1 Task code ★☆☆ [~1 week]

3.2 Training ★☆☆⚠ [~2 week]

3.3 Analysis code ★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆

1.1 Code logic for new environment [~1 week]

1.2 Check everything works as expected (unit testing)
[~1 week]

1.3 Bugs? [~1 week]

2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★☆☆⚠

[1 week – 1 month]

3. Experiments

3.1 Task code ★☆☆ [~1 week]

3.2 Training ★☆☆⚠ [~2 week]

3.3 Analysis code ★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★★☆

1.1 Code logic for new environment [~1 week]

1.2 Check everything works as expected (unit testing)
[~1 week]

1.3 Bugs? [~1 week]

2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★★⚠

[1 week – 1 month]

3. Experiments

3.1 Task code ★☆☆ [~1 week]

3.2 Training ★☆☆ [~2 week]

3.3 Analysis code ★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆
 - 1.1 Code logic for new environment [~1 week]
 - 1.2 Check everything works as expected (unit testing)
[~1 week]
 - 1.3 Bugs? [~1 week]
2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★★⚠
[1 week – 1 month]
3. Experiments
 - 3.1 Task code ★☆☆ [~1 week]
 - 3.2 Training ★★☆☆⚠ [~2 week]
 - 3.3 Analysis code ★★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆
 - 1.1 Code logic for new environment [~1 week]
 - 1.2 Check everything works as expected (unit testing)
[~1 week]
 - 1.3 Bugs? [~1 week]
2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★★⚠
[1 week – 1 month]
3. Experiments
 - 3.1 Task code ★☆☆ [~1 week]
 - 3.2 Training ★★☆☆⚠ [~2 week]
 - 3.3 Analysis code ★★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆
 - 1.1 Code logic for new environment [~1 week]
 - 1.2 Check everything works as expected (unit testing)
[~1 week]
 - 1.3 Bugs? [~1 week]
2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★★⚠
[1 week – 1 month]
3. Experiments
 - 3.1 Task code ★☆☆ [~1 week]
 - 3.2 Training ★★★⚠ [~2 week]
 - 3.3 Analysis code ★☆☆ [~2 week]

Milestones/how to get there

1. Rewrite the environment(s) ★★☆☆
 - 1.1 Code logic for new environment [~1 week]
 - 1.2 Check everything works as expected (unit testing)
[~1 week]
 - 1.3 Bugs? [~1 week]
2. Baseline training on new environment (convergence, hyperparameter tweaking, etc.) ★★★⚠
[1 week – 1 month]
3. Experiments
 - 3.1 Task code ★☆☆ [~1 week]
 - 3.2 Training ★★★⚠ [~2 week]
 - 3.3 Analysis code ★★★ [~2 week]

Thanks!