DRL project update

Cartesian/polar duplicated coordinates experiment

Andrea Pierré

January 21, 2025

Outline

1. Current status

2. How to get insights at what the network learn?

Outline

1. Current status

2. How to get insights at what the network learn?

Current status

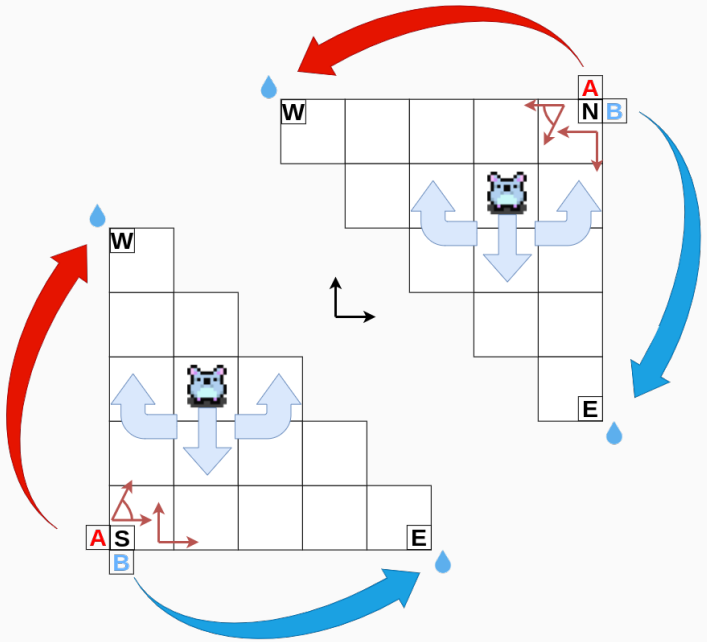

- Environment rewrite: done
- Training: ~4.5 hours to train 30 agents on both tasks on Oscar
- Analysis: WIP

State space & network architecture

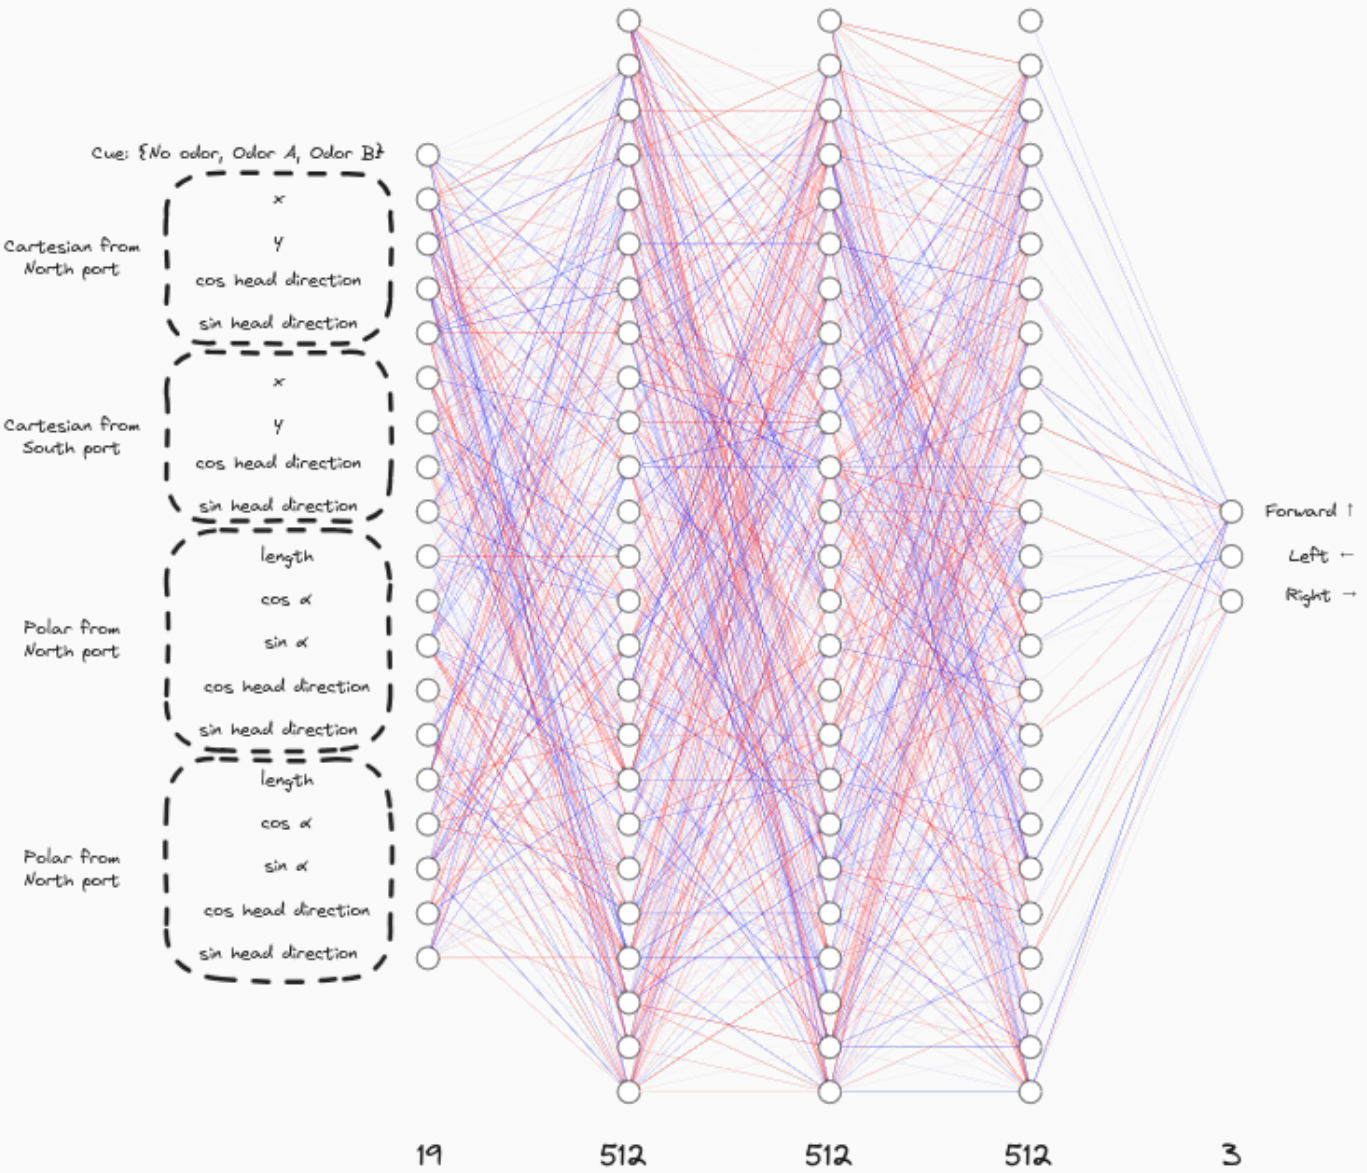

Training

East/West

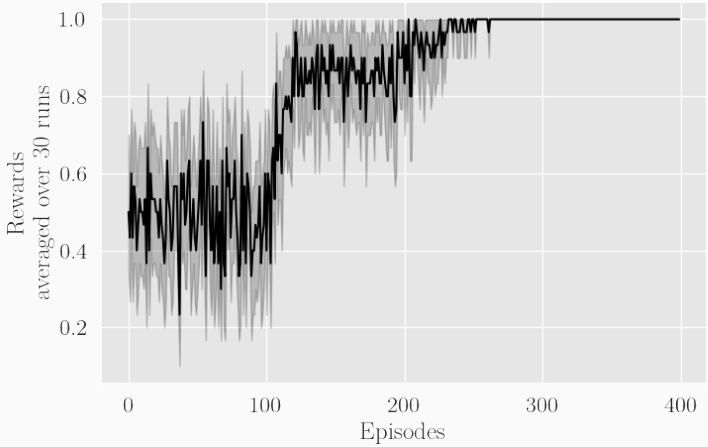

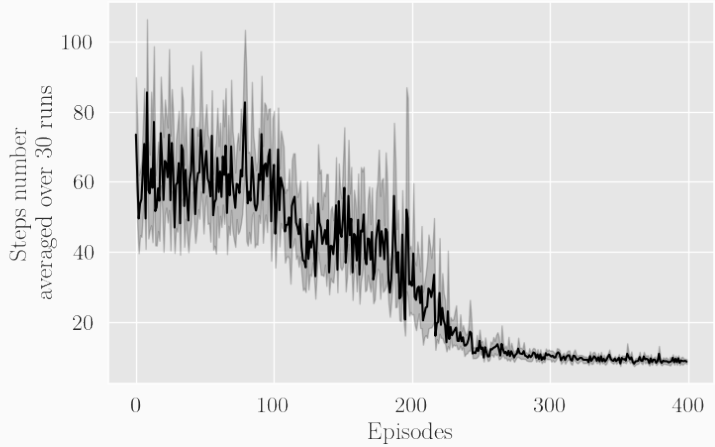

Left/Right

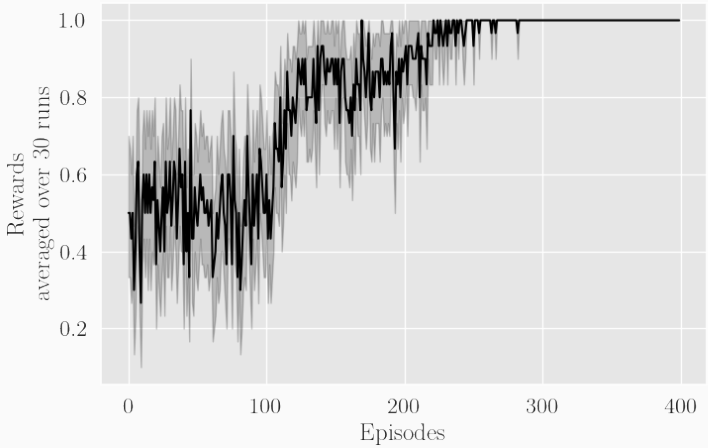

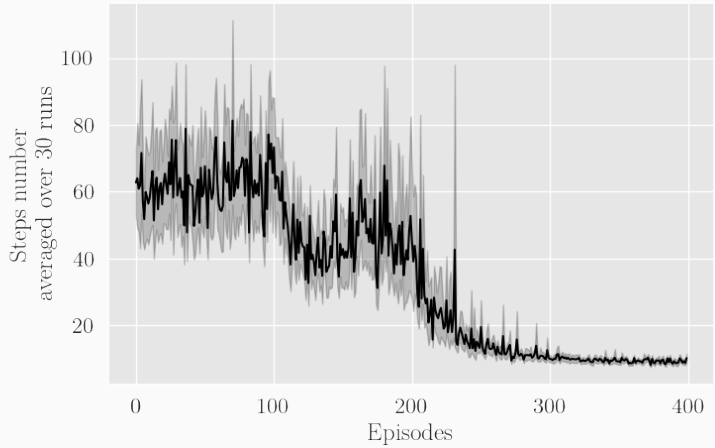

Training checks - East/West

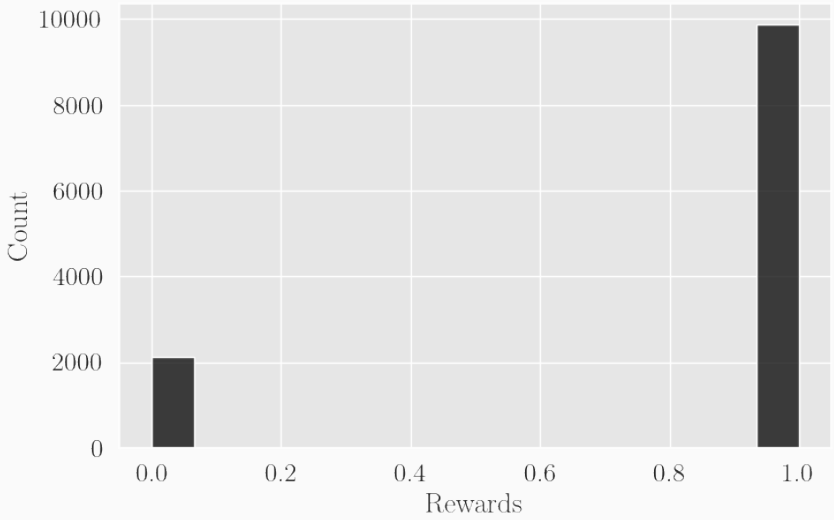

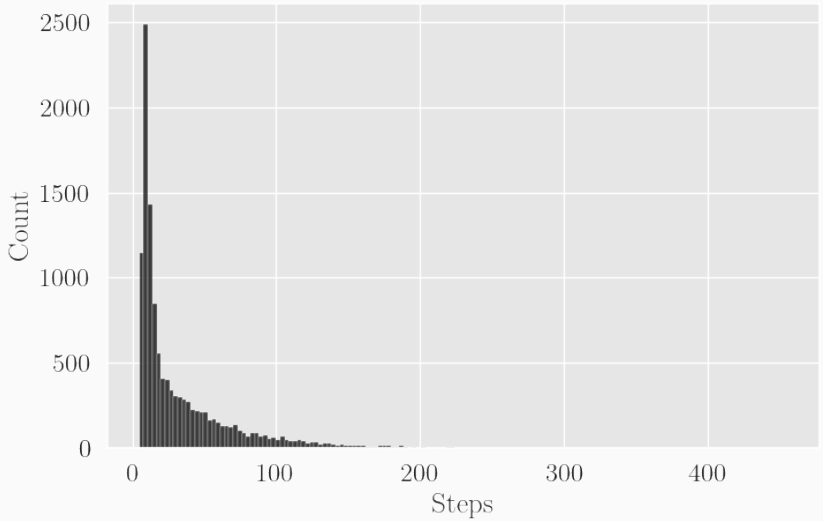

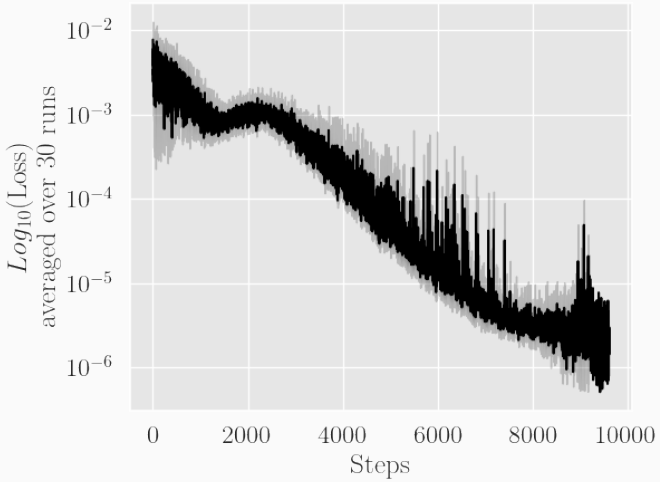

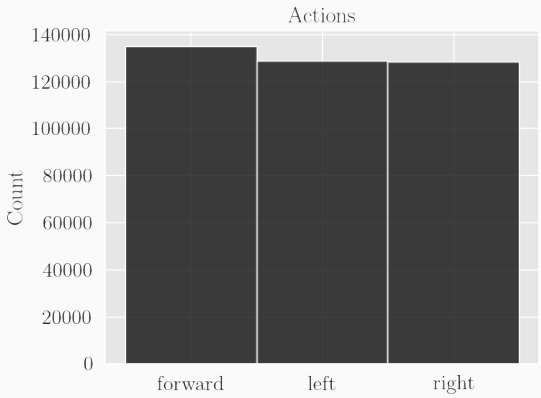

Training checks - Left/Right

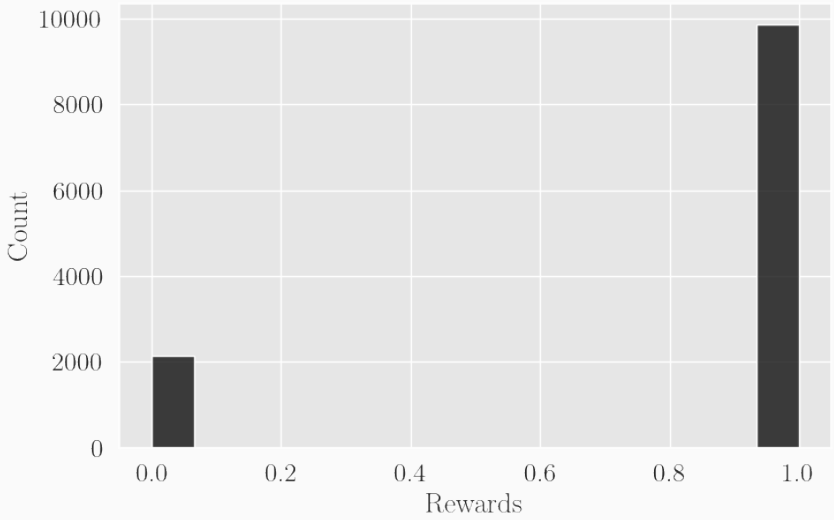

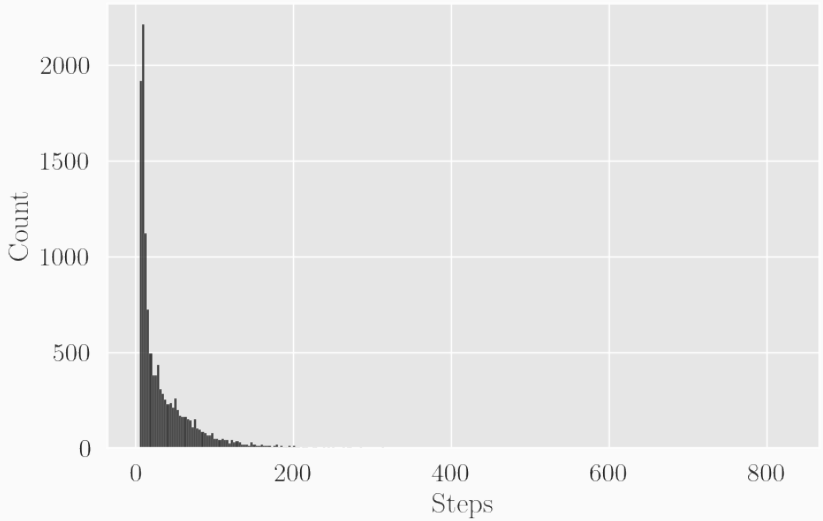

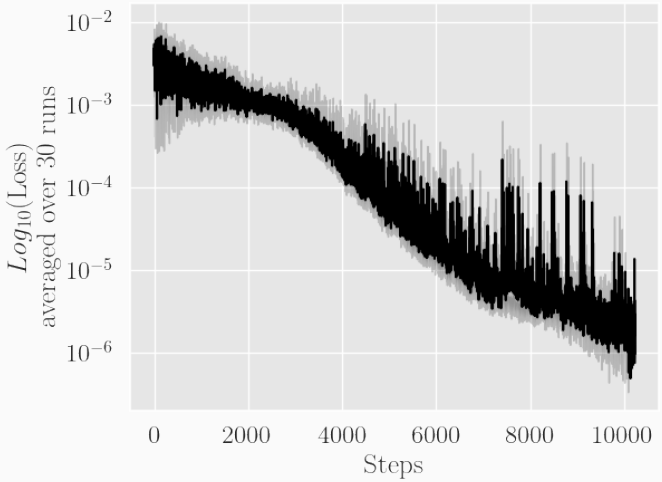

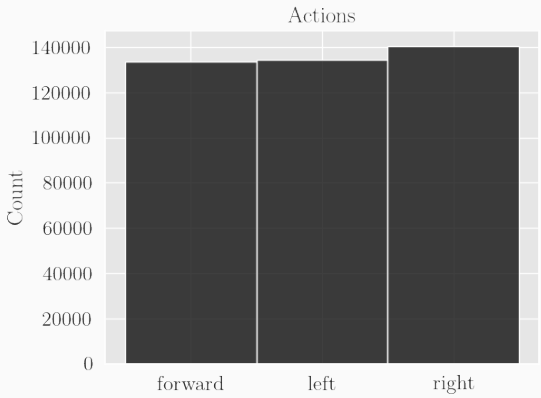

Agent behavior

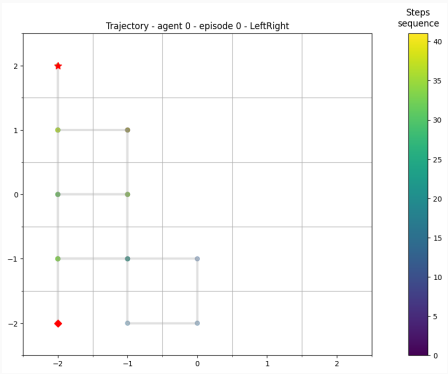

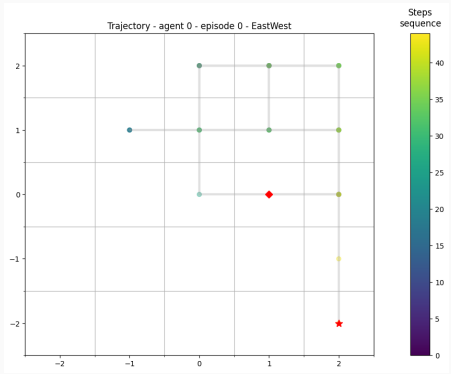

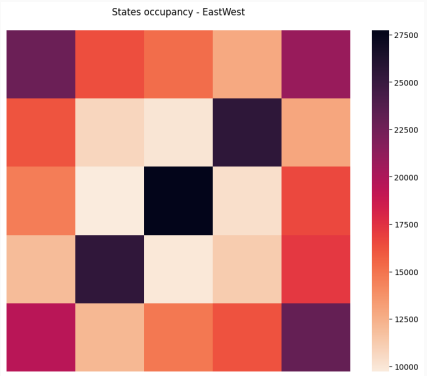

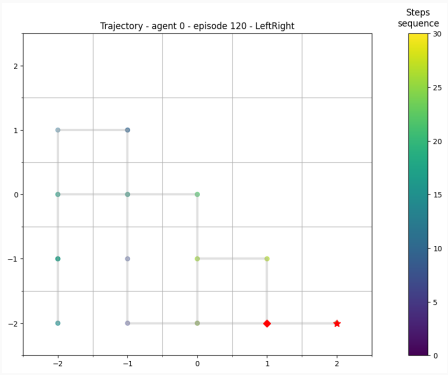

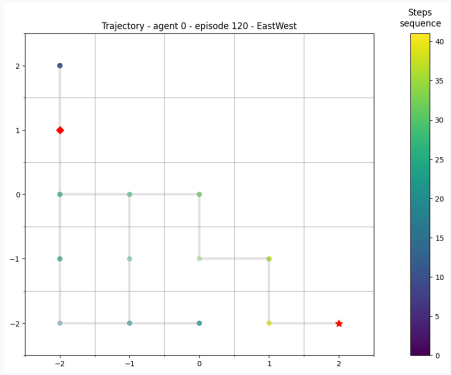

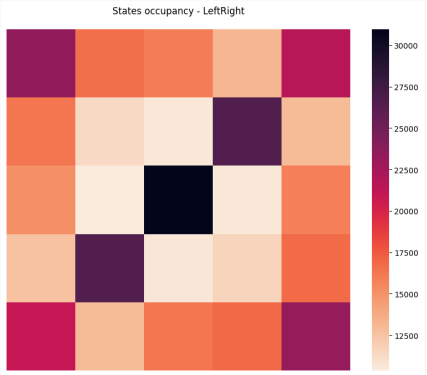

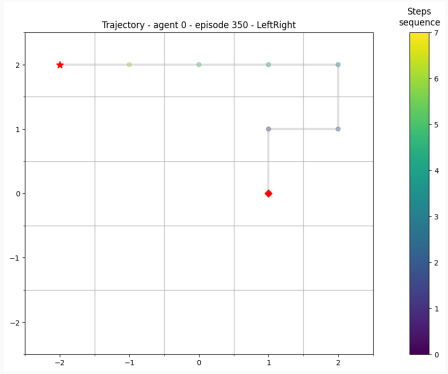

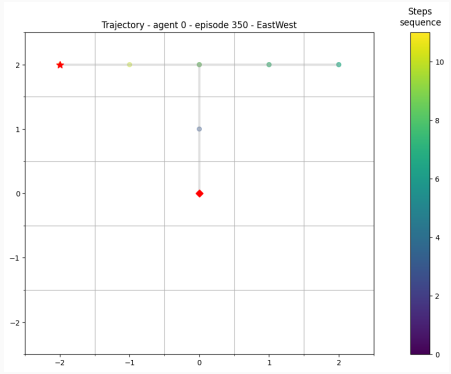

Outline

1. Current status

2. How to get insights at what the network learn?

Activations learned - East/West

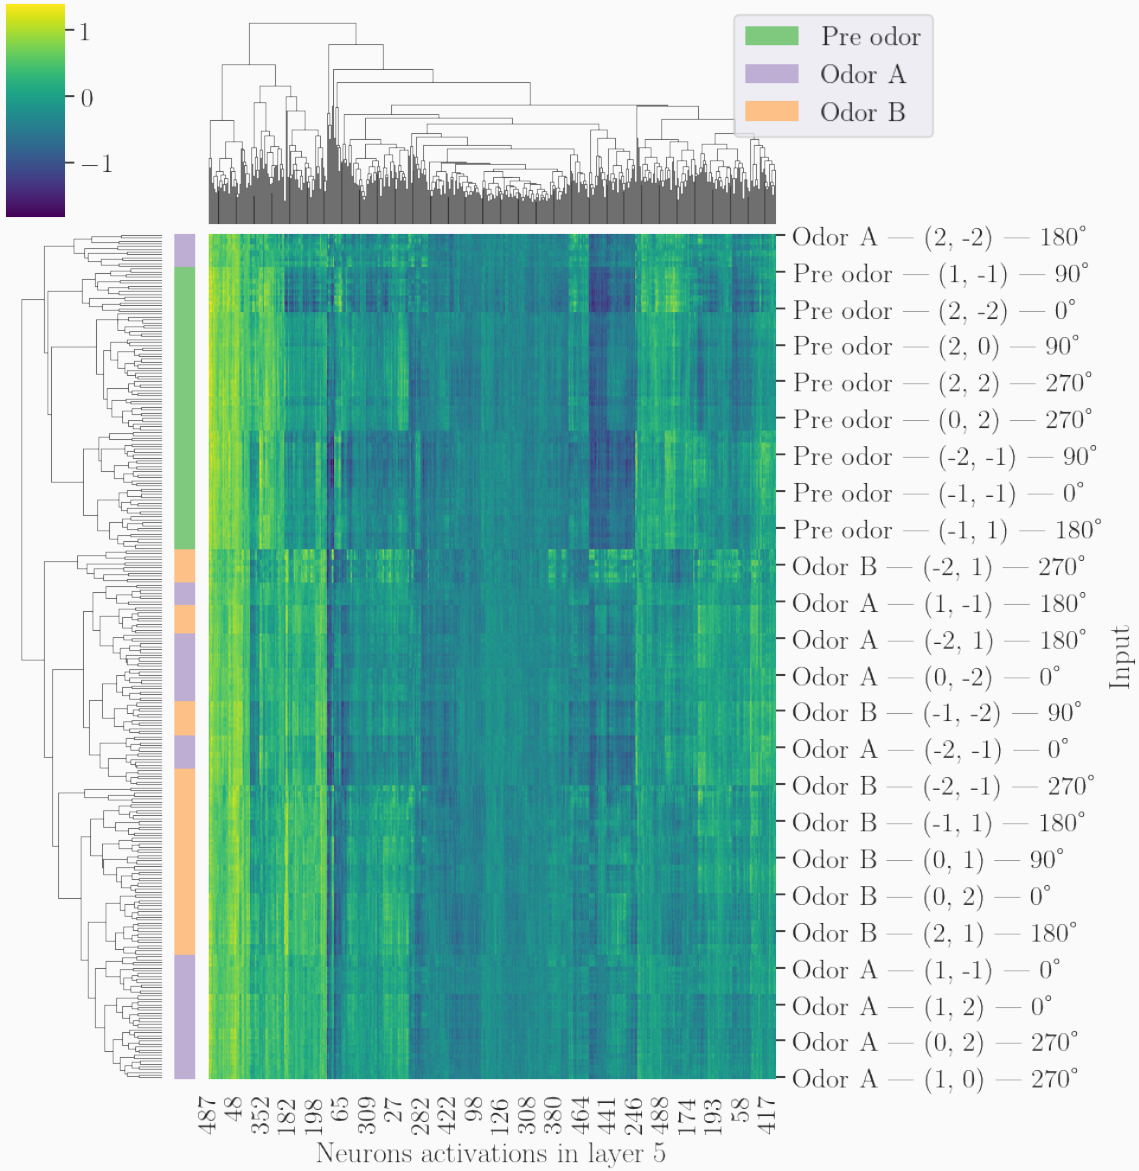

Activations learned - Left/Right

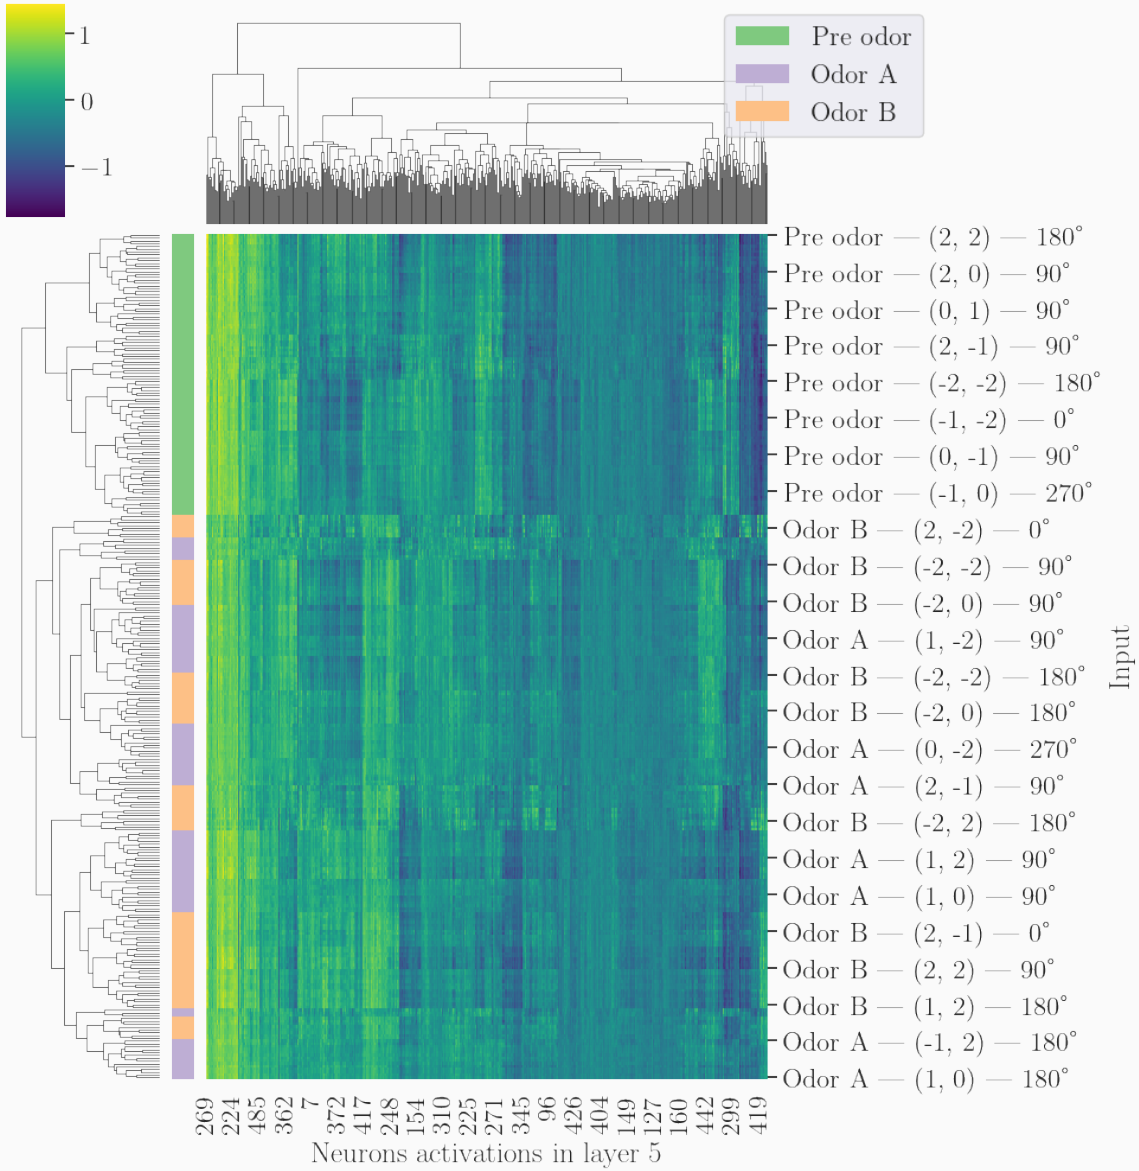

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
  - East/west task:
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced → the agent can solve the task
    - With the polar inputs silenced → the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced → the agent can solve the task
    - With the polar inputs silenced → the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced → the agent can solve the task
    - With the polar inputs silenced → the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced → the agent can solve the task
    - With the polar inputs silenced → the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
- Any other approach we could use?

Use the behavior as proxy

- Silence the Cartesian/polar part of the input on a trained agent and look at how the agent behaves (x4 experiments)
- Expectation:
  - Left/right task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
  - East/west task:
    - With the Cartesian inputs silenced $\rightarrow$ the agent can solve the task
    - With the polar inputs silenced $\rightarrow$ the agent struggle to solve the task
- Any other approach we could use?

Neural representations?

2. How the constraints of the task impact the representations learned?

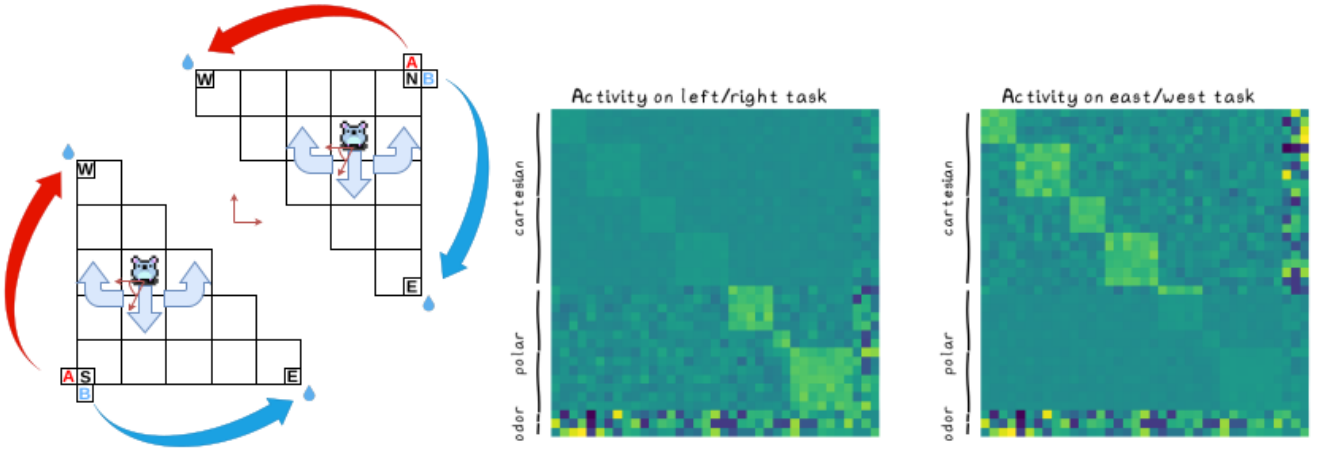

(a) Experiment with duplicated spatial information input.

(b) Expected neuronal activity on the left/right task (left) and on the east/west task (right).

Figure 8. What does the network learn if it gets both Cartesian and polar information as input?
On the left/right task → we expect the activity to be close to zero on the Cartesian representation.
On the east/west task → we expect the activity to be close to zero on the polar representation.

- Need for some causal framework?